



TensorFlow Tutorial

CS 419(M) : Introduction to Machine Learning

Prof. Sunita Sarawagi

Divyansh Pareek and Karan Taneja



TensorFlow : What is it ?

- Open source software library
- Intensely optimized for fast Numerical computations
- Widely used for Machine Learning / Deep Learning



Getting Started

- Installation
 - OS dependent instructions available on <https://www.tensorflow.org/install/>
 - Install CPU version (instead of GPU) if you're using a laptop
- Google Colab
 - Jupyter notebook environment
 - Good for getting off-the-ground
 - Go to <https://colab.research.google.com> and start typing!

We'll not cover the basics of python, we assume you are somewhat familiar with it.



Colab : Create a notebook

- Go to colab
- Create a python3 (Jupyter) notebook
- Go to *Edit > Notebook Settings*
 - Change Hardware Accelerator from None to GPU
- Click Connect to get connected to a machine in Google's datacenter
- Start typing!

But first ... Tensor : What is it (in TF) ?

- Generalization of vectors and matrices to potentially higher dimensions
- Everything is a tensor!
 - A scalar is a tensor of dimension 0
 - A vector is a tensor of dimension 1
 - A matrix is a tensor of dimension 2

't'
'e'
'n'
's'
'o'
'r'

Tensor of dimension[1]

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

Tensor of dimensions[2]

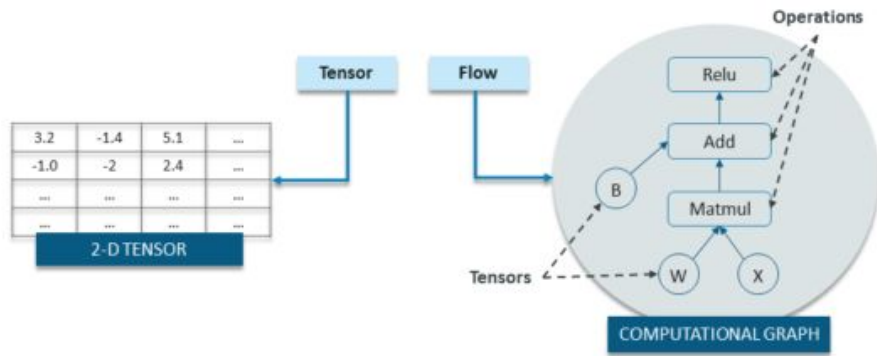
2	1	2	8	1	8
2	8	5	9	0	4
2	3	3	0	2	8
7	7	1	3	5	2

Tensor of dimensions[3]

- Say you have N images, each of size H x W - represent that as a tensor of dimension N x H x W
 - Each entry would be a pixel value, say an *int* (base datatype) in the range [0, 256)
 - What about N videos, each of H x W spatial dimension and T units in time dimension ?
- Internally, TensorFlow represents tensors as n-dimensional arrays of base datatypes (*tf.int32*, *tf.float64*, etc)

TensorFlow = Tensor + Flow (Duh!)

- TensorFlow is quite literally the flow of tensors



- The idea of a **computational graph** is central to TensorFlow
 - We'll see this in more detail shortly

First TF Program : Adding two tensors

```
import tensorflow as tf

# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
b = tf.constant([1.0,2.0,4.0])

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b # operator + is overloaded for tensors

with tf.Session() as sess:
    c_value = sess.run(c) # Could also have done : c_value = c.eval()
    print("Value of Node c : ", c_value)
    print("Attributes of Value of Node c : ", "Type:", type(c_value), " Shape:", c_value.shape)
    print("----")
    print("Node c itself : ", c)
    print("Shape of Node c : ", c.shape)
    print("Dtype of Node c : ", c.dtype)
```

```
Value of Node c : [2. 4. 7.]
Attributes of Value of Node c : Type: <class 'numpy.ndarray'> Shape: (3,)
----
Node c itself : Tensor("add_1:0", shape=(3,), dtype=float32)
Shape of Node c : (3,)
Dtype of Node c : <dtype: 'float32'>
```

← Output of sess.run(c) is a np.ndarray

← c itself is a tf.tensor

Computational Graph Model

- The fragment of the code before the initialization of `tf.Session()` was this


```
# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
b = tf.constant([1.0,2.0,4.0])

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b

with tf.Session() as sess:
```

- This defines a computation graph as follows



The arrows represent that the definition of `c` requires `a` & `b` (since `c := a + b`).

This computational graph is independent of the values that `a` & `b` hold (although here they are constants). `tf` internally builds (and stores) this graph.



Computational Graph Model : `tf.Session()`

- Writing code in tf => defining such computation graphs (for whatever operations we want)
 - UNLIKE normal python/numpy, this ONLY defines the computation graph
 - ie, it just defines all the operations on all the various tensors that we want
 - Does NOT carry out any computation, just creates and stores a computation graph
-
- Then, instantiate a `tf.Session()`
 - Call `sess.run(tens)`
 - Here `sess` is the name of the `tf.Session()` object we created
 - And `tens` is any tensor whose value we want to compute
 - This call makes tf run a forward pass on the graph that it created, using actual values
 - This is the step where the computation actually happens

TensorFlow computations define a computation graph that has no numerical value until evaluated!



Computational Graph Model

- One of the central ideas of tensorflow
- “TensorFlow programs are usually structured into a construction phase, that assembles a graph, and an execution phase that uses a session to execute ops in the graph.” - TensorFlow docs
- “A Session object encapsulates the environment in which Operation objects are executed, and Tensor objects are evaluated.” - TensorFlow Docs
- All computations add nodes to global default graph (docs)
 - `tf.Graph` is the class handling such computation graphs
 - Lots of operations to modify and play around with computational graphs
 - To handle the default graph: `g = tf.get_default_graph()` ... then modify `g`
 - You can create multiple computation graphs
 - This will just be multiple instances of `tf.Graph`

Computational Graph Model

```
import tensorflow as tf

# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
b = tf.constant([1.0,2.0,4.0])

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b # operator + is overloaded for tensors

with tf.Session() as sess:
    c_value = sess.run(c) # Could also have done : c_value = c.eval()
    print("Value of Node c : ", c_value)
    print("Attributes of Value of Node c : ", "Type:", type(c_value), " Shape:", c_value.shape)
    print("----")
    print("Node c itself : ", c)
    print("Shape of Node c : ", c.shape)
    print("Dtype of Node c : ", c.dtype)
```

Define a computation graph. At this point we can't see real value.

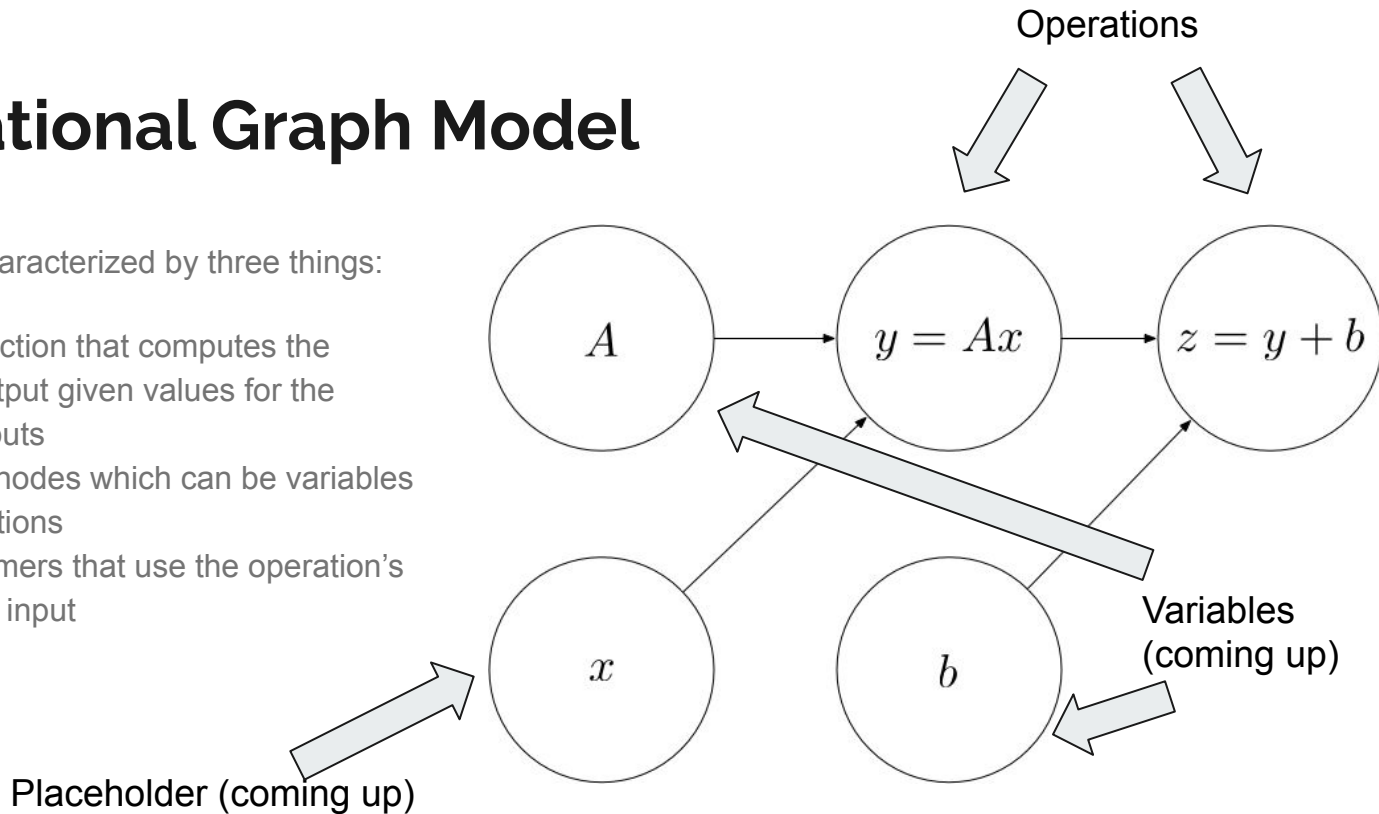
Evaluate pre-defined comp graph. Graph eval must be preceded by a `tf.Session()` or `tf.InteractiveSession()` creation.

```
Value of Node c : [2. 4. 7.]
Attributes of Value of Node c : Type: <class 'numpy.ndarray'> Shape: (3,)
----
Node c itself : Tensor("add_1:0", shape=(3,), dtype=float32)
Shape of Node c : (3,)
Dtype of Node c : <dtype: 'float32'>
```

Computational Graph Model

Every operation is characterized by three things:

- A compute function that computes the operation's output given values for the operation's inputs
- A list of input_nodes which can be variables or other operations
- A list of consumers that use the operation's output as their input



Playing with Add : explicit shape def

```
[5] import tensorflow as tf

# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,1) (specified) | datatype = tf.float32 (inferred)
b = tf.constant([1.0,2.0,4.0], shape=(3,1)) # NO ERROR, but wasn't what we intended -- does elementwise addition
# interpreted as [1,2,3] + [[1],[2],[4]] -- elementwise addition done on this

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b # operator + is overloaded for tensors

with tf.Session() as sess:
    c_value = sess.run(c) # Could also have done : c_value = c.eval()
    print("Value of Node c : ", c_value)
    print("Attributes of Value of Node c : ", "Type:", type(c_value), " Shape:", c_value.shape)
    print("---- --- ----")
    print("Node c itself : ", c)
    print("Shape of Node c : ", c.shape)
    print("Dtype of Node c : ", c.dtype)
```

```
> Value of Node c :  [[2. 3. 4.]
 [3. 4. 5.]
 [5. 6. 7.]]
Attributes of Value of Node c :  Type: <class 'numpy.ndarray'>  Shape: (3, 3)
---- --- ----
Node c itself :  Tensor("add_3:0", shape=(3, 3), dtype=float32)
Shape of Node c :  (3, 3)
Dtype of Node c :  <dtype: 'float32'>
```

Playing with Add : explicit dtype def

```
[7] import tensorflow as tf

# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,) (specified) | datatype = tf.float64 (specified)
b = tf.constant([1.0,2.0,4.0], shape=(3,), dtype=tf.float64) # same as the original

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b # operator + is overloaded for tensors

with tf.Session() as sess:
    c_value = sess.run(c) # Could also have done : c_value = c.eval()
    print("Value of Node c : ", c_value)
    print("Attributes of Value of Node c : ", "Type:", type(c_value), " Shape:", c_value.shape)
    print("----")
    print("Node c itself : ", c)
    print("Shape of Node c : ", c.shape)
    print("Dtype of Node c : ", c.dtype)
```



```
-----
ValueError                                Traceback (most recent call last)
/usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py in _apply_op_helper(self, op_type_name, name, **keywords)
    509         as_ref=input_arg.is_ref,
--> 510         preferred_dtype=default_dtype)
    511     except TypeError as err:

/usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/ops.py in internal_convert_to_tensor(value, dtype, name, as_ref, preferred_dtype)
    1002         if not is None:
```

Playing with Add : using tf.add and node

```
[1] import tensorflow as tf

# Declare a tensor with : constant value = [1,2,3] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
a = tf.constant([1.0,2.0,3.0])

# Declare a tensor with : constant value = [1,2,4] | dimension = (3,) (inferred) | datatype = tf.float32 (inferred)
b = tf.constant([1.0,2.0,4.0])

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = tf.add(a, b, name='MyAdd') # using tf.add instead of overloaded operator +

with tf.Session() as sess:
    c_value = sess.run(c) # Could also have done : c_value = c.eval()
    print("Value of Node c : ", c_value)
    print("Attributes of Value of Node c : ", "Type:", type(c_value), " Shape:", c_value.shape)
    print("---- --- ---")
    print("Node c itself : ", c)
    print("Shape of Node c : ", c.shape)
    print("Dtype of Node c : ", c.dtype)
```

```
↳ Value of Node c : [2. 4. 7.]
Attributes of Value of Node c : Type: <class 'numpy.ndarray'> Shape: (3,)
---- --- ---
Node c itself : Tensor("MyAdd:0", shape=(3,), dtype=float32)
Shape of Node c : (3,)
Dtype of Node c : <dtype: 'float32'>
```



Non-constant data : Placeholders and feed dictionaries

- What we did was pretty trivial, adding two constant vectors
 - Let's say we want to still add two vectors, but now they are not constants
 - How do we handle non-constant ? *tf.placeholder*!
-
- *tf.placeholder* defines dummy nodes that provide entry points for data to computational graph
 - A *feed_dict* is a python dictionary mapping from *tf.placeholder* vars (or their names) to data (numpy arrays, lists, etc.)

Placeholders and feed_dict : Example

```
import tensorflow as tf

# Declare a placeholder for a : actual data provided at runtime
a = tf.placeholder(tf.float32)

# Declare a placeholder for b : actual data provided at runtime
b = tf.placeholder(tf.float32)

# Add a & b
# c is a Tensor - the result of addition of tensors a & b
c = a + b

with tf.Session() as sess:
    c_value = sess.run(c, feed_dict={a:[1.0,2.0,3.0], b:[1.0,2.0,4.0]}) # can give anything here ... in particular, read a dataset and give that here
    print("Value of Node c : ", c_value)
```

Value of Node c : [2. 4. 7.]

Analogy: Can define the function $f(x, y) = x^2 + y$ without knowing value of x or y .
 x , y are placeholders for the actual values.

Will give error if mismatch during computation (ie, forward pass on graph)

Activities

Google Chrome

TensorFlow_Tutor

Examples for tf tu

tf_demo.ipynb - C

tf.placeholder

Tensors

Tensor

TensorFlow Tutor

CS224d-Lecture7

Google TensorFlo

TensorFlow Tutor

@divyansh

Secure | https://www.tensorflow.org/api_docs/python/tf/placeholder

Apps Gmail CSE, IIT Bombay IIT Bombay Moo Deep RL Bootca CS 236: Deep Ge CS 228: Probabili Firefox Chrome Suits EM Alg Music CSRankings: Con Other bookmarks

TensorFlow™

Install Develop Community API r1.10 Ecosystem

Search

GITHUB

PYTHON

JAVASCRIPT

C++

JAVA

SERVING

SWIFT

MORE...

OptimizerOptions

op_scope

orthogonal_initializer

pad

PaddingFIFOQueue

parallel_stack

parse_example

parse_single_example

parse_single_sequence_example

parse_tensor

placeholder

placeholder_with_default

polygamma

pow

Print

PriorityQueue

py_func

qr

quantize

quantized_concat

quantize_v2

QueueBase

RandomShuffleQueue

random_crop

random_gamma

random_normal

random_normal_initializer

random_poisson

random_shuffle

random_uniform

random_uniform_initializer

range

rank

tf.placeholder

☆☆☆☆☆

```
tf.placeholder(  
    dtype,  
    shape=None,  
    name=None  
)
```

Defined in [tensorflow/python/ops/array_ops.py](#).

See the guides: [Inputs and Readers > Placeholders](#), [Reading data > Feeding](#)

Inserts a placeholder for a tensor that will be always fed.

Important: This tensor will produce an error if evaluated. Its value must be fed using the `feed_dict` optional argument to `Session.run()`, `Tensor.eval()`, or `Operation.run()`.

For example:

```
x = tf.placeholder(tf.float32, shape=(1024, 1024))  
y = tf.matmul(x, x)  
  
with tf.Session() as sess:  
    print(sess.run(y)) # ERROR: will fail because x was not fed.  
  
    rand_array = np.random.rand(1024, 1024)  
    print(sess.run(y, feed_dict={x: rand_array})) # Will succeed.
```

Args:

- `dtype`: The type of elements in the tensor to be fed.
- `shape`: The shape of the tensor to be fed (optional). If the shape is not specified, you can feed a tensor of any shape.
- `name`: A name for the operation (optional).

Will give error if mismatch during runtime ...



Next Major piece of the puzzle : `tf.Variable()`

- We've seen
 - Computational graph model
 - Placeholders - special tensors that can be fed with input data
- With these, we can only do what all we could do in numpy
- The next major piece is `tf.Variable()`
 - “A variable maintains state in the graph across calls to `run()`” - Tensorflow Docs
 - “When you train a model you use variables to hold and update parameters. Variables are in-memory buffers containing tensors” - TensorFlow Docs

Variables

Why `tf.constant` but `tf.Variable` and not `tf.variable`?

`tf.Variable` is a class, but `tf.constant` is an op

`tf.Variable` holds several ops:

`x = tf.Variable(...)`

`x.initializer` # init op

`x.value()` # read op

`x.assign(...)` # write op

- Variable allows you to add such parameters or node to the graph that are trainable
 - i.e. the value can be modified over the period of a time
- “The `Variable()` constructor requires an initial value for the variable, which can be a Tensor of any type and shape. The initial value defines the type and shape of the variable. After construction, the type and shape of the variable are fixed.” - TensorFlow Docs
- Example

```
import tensorflow as tf
```

```
# Create a variable.
```

```
w = tf.Variable(<initial-value>, name=<optional-name>)
```

```
# Use the variable in the graph like any Tensor.
```

```
y = tf.matmul(w, ...another variable or tensor...)
```

```
# The overloaded operators are available too.
```

```
z = tf.sigmoid(w + y)
```

```
# Assign a new value to the variable with `assign()` or a related method.
```

```
w.assign(w + 1.0)
```

```
w.assign_add(1.0)
```

Variables : Need to be initialized

```
import tensorflow as tf

# Declare a variable
w = tf.Variable(tf.zeros((2,2)), name="weight") # initialized by the 2 x 2 matrix of zeros

with tf.Session() as sess:
    sess.run(w)
```

```
-----
FailedPreconditionError                                Traceback (most recent call last)
/usr/local/lib/python3.6/dist-packages/tensorflow/python/client/session.py in _do_call(self, fn, *args)
    1277         try:
-> 1278             return fn(*args)
```

```
[8] import tensorflow as tf

# Declare a variable
w = tf.Variable(tf.zeros((2,3)), name="weight") # initialized by the 2 x 2 matrix of zeros

with tf.Session() as sess:
    sess.run(w.initializer)
    w_value = sess.run(w)
    print("Value of w : ", w_value)
```

```
Value of w :  [[0. 0. 0.]
 [0. 0. 0.]]
```



Variables : initialization

The easiest way is initializing all variables at once:

```
init = tf.global_variables_initializer()
```

with *tf.Session()* as *sess*:

```
sess.run(init)
```

Initialize only a subset of variables:

```
init_ab = tf.variables_initializer([a, b], name="init_ab")
```

with *tf.Session()* as *sess*:

```
sess.run(init_ab)
```

Initialize a single variable:

```
W = tf.Variable(tf.zeros([784,10]))
```

with *tf.Session()* as *sess*:

```
sess.run(W.initializer)
```



Variables : Updation

- Since variables maintain state across runs, we naturally want them to update across runs
- How do we update them ?

```
W = tf.Variable(10)
W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    print W.eval() # >> 10, not 100
```

```
W = tf.Variable(10)
update_op = W.assign(100)
with tf.Session() as sess:
    sess.run(W.initializer)
    sess.run(update_op)
    print W.eval() # >> 100
```

- Note that we defined an operation for the variable updation
 - And we ran that op using sess.run() to update the variable



Variables across sessions

```
W = tf.Variable(10)
sess1 = tf.Session()
sess2 = tf.Session()
sess1.run(W.initializer)
sess2.run(W.initializer)
print sess1.run(W.assign_add(10)) # >> 20
print sess2.run(W.assign_sub(2)) # >> 8
print sess1.run(W.assign_add(100)) # >> 120
print sess2.run(W.assign_sub(50)) # >> -42
sess1.close()
sess2.close()
```




Let's do a proper example!

- [Example 1](#) (Slide 27)
- [Example 2](#) (Slide 31)



More awesome(!) things about tf

What we've covered are just the very basics. But this should enable you to read & understand tf code more easily.

tf.device -- multiple devices

tf.stop_gradient -- stop gradient somewhere in computational graph

tf.Print -- a node that prints out its contents at runtime

tf.trainable_variables -- make some variables trainable

tf.get_variable

tf.layers.* (conv2d, dropout, pooling, etc)

tf.nn.* (relu, etc)

tf.reduce_*

tf.train.* (optimizers, etc)

Lots and lots and lots more!

Selling points of tf:

1. Autodiff (nowadays all libraries have this)
2. Highly optimised for both graph creation/storage and also graph operations (this is the real power of tf)
3. Tensorboard - powerful visualisation tool

Many others...



References

https://www.tensorflow.org/api_docs/python/

<https://www.edureka.co/blog/tensorflow-tutorial/>

<https://cs224d.stanford.edu/lectures/CS224d-Lecture7.pdf>

https://www.slideshare.net/tw_dsconf/tensorflow-tutorial

<https://www.slideshare.net/nmhkahn/tensorflow-tutorial-71896086>

https://web.stanford.edu/class/cs20si/2017/lectures/slides_01.pdf



Thank You!